

```
In [4]: import pandas as pd

df = pd.read_csv(r"C:\Users\Mummy Mo\Downloads\HNG Data set\marketing_campaign_dataset.csv")

# Display basic information about the dataset
df_info = df.info()
df_head = df.head()

df_info, df_head
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200005 entries, 0 to 200004
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Campaign_ID           200005 non-null  int64
1   Company               200005 non-null  object
2   Campaign_Type         200005 non-null  object
3   Target_Audience      200005 non-null  object
4   Duration              200005 non-null  object
5   Channel_Used          200005 non-null  object
6   Conversion_Rate       200005 non-null  float64
7   Acquisition_Cost      200005 non-null  object
8   ROI                  200005 non-null  float64
9   Location              200005 non-null  object
10  Date                  200005 non-null  object
11  Clicks                200005 non-null  int64
12  Impressions           200005 non-null  int64
13  Engagement_Score      200005 non-null  int64
14  Customer_Segment      200005 non-null  object
dtypes: float64(2), int64(4), object(9)
memory usage: 22.9+ MB
```

```
Out[4]: (None,
         Campaign_ID      Company Campaign_Type Target_Audience Duration \
0          1  Innovate Industries      Email      Men 18-24  30 days
1          2    NexGen Systems      Email      Women 35-44  60 days
2          3  Alpha Innovations  Influencer      Men 25-34  30 days
3          4  DataTech Solutions  Display      All Ages  60 days
4          5    NexGen Systems      Email      Men 25-34  15 days

         Channel_Used Conversion_Rate Acquisition_Cost  ROI      Location \
0   Google Ads         0.04      $16,174.00  6.29      Chicago
1   Google Ads         0.12      $11,566.00  5.61      New York
2   YouTube           0.07      $10,200.00  7.18  Los Angeles
3   YouTube           0.11      $12,724.00  5.55      Miami
4   YouTube           0.05      $16,452.00  6.50  Los Angeles

         Date Clicks Impressions Engagement_Score Customer_Segment
0  01/01/2021   506      1922              6  Health & Wellness
1  01/02/2021   116      7523              7      Fashionistas
2  01/03/2021   584      7698              1  Outdoor Adventurers
3  01/04/2021   217      1820              7  Health & Wellness
4  01/05/2021   379      4201              3  Health & Wellness )
```

```
In [10]: # Convert 'Date' to datetime format
df['Date'] = pd.to_datetime(df['Date'], format="%d/%m/%Y")

# Convert 'Acquisition_Cost' to numeric (removing '$' and commas)
df['Acquisition_Cost'] = df['Acquisition_Cost'].replace(r'[\$,]', '', regex=True).astype(float)
```

```
In [12]: # Check for missing values
missing_values = df.isnull().sum()

# Generate summary statistics for numerical columns
summary_stats = df.describe()

missing_values, summary_stats
```

```

Out[12]: (Campaign_ID          0
          Company              0
          Campaign_Type        0
          Target_Audience      0
          Duration              0
          Channel_Used          0
          Conversion_Rate       0
          Acquisition_Cost      0
          ROI                  0
          Location              0
          Date                  0
          Clicks                0
          Impressions           0
          Engagement_Score      0
          Customer_Segment      0
          dtype: int64,

          Campaign_ID  Conversion_Rate  Acquisition_Cost  ROI \
count  200005.000000    200005.000000    200005.000000  200005.000000
mean   100003.000000         0.080069      12504.441794      5.002416
min     1.000000          0.010000       5000.000000      2.000000
25%     50002.000000         0.050000       8740.000000      3.500000
50%     100003.000000        0.080000      12497.000000      5.010000
75%     150004.000000        0.120000      16264.000000      6.510000
max     200005.000000        0.150000      20000.000000      8.000000
std      57736.614632         0.040602       4337.663210      1.734485

          Date              Clicks  Impressions \
count              200005    200005.000000  200005.000000
mean  2021-07-01 23:37:44.289392896      549.774591    5507.307107
min    2021-01-01 00:00:00          100.000000    1000.000000
25%    2021-04-02 00:00:00          325.000000    3266.000000
50%    2021-07-02 00:00:00          550.000000    5518.000000
75%    2021-10-01 00:00:00          775.000000    7753.000000
max    2021-12-31 00:00:00         1000.000000   10000.000000
std              NaN          260.019354    2596.863794

          Engagement_Score
count    200005.000000
mean         5.494673
min          1.000000
25%          3.000000
50%          5.000000
75%          8.000000
max         10.000000
std          2.872593 )

```

```

In [14]: # Check for any remaining null values in 'Date' after conversion
date_null_count = df['Date'].isnull().sum()

# Display rows where 'Date' could not be converted (if any)
invalid_dates = df[df['Date'].isnull()].head()

date_null_count, invalid_dates

```

```

Out[14]: (0,
          Empty DataFrame
          Columns: [Campaign_ID, Company, Campaign_Type, Target_Audience, Duration, Channel_Used, Conversion_Rate, Acquisition_Cost, ROI, Location, Date, Clicks, Impressions, Engagement_Score, Customer_Segment]
          Index: [])

```

```

In [16]: # Print column names to check for typos
print(df.columns)

Index(['Campaign_ID', 'Company', 'Campaign_Type', 'Target_Audience',
       'Duration', 'Channel_Used', 'Conversion_Rate', 'Acquisition_Cost',
       'ROI', 'Location', 'Date', 'Clicks', 'Impressions', 'Engagement_Score',
       'Customer_Segment'],
      dtype='object')

```

```

In [18]: # Check for missing values again after cleaning
missing_values = df.isnull().sum()

# Generate summary statistics for numerical columns
summary_stats = df.describe()

missing_values, summary_stats

```

```

Out[18]: (Campaign_ID          0
          Company            0
          Campaign_Type      0
          Target_Audience    0
          Duration            0
          Channel_Used        0
          Conversion_Rate     0
          Acquisition_Cost    0
          ROI                 0
          Location            0
          Date                0
          Clicks              0
          Impressions         0
          Engagement_Score    0
          Customer_Segment    0
          dtype: int64,

          Campaign_ID  Conversion_Rate  Acquisition_Cost  ROI \
count  200005.000000    200005.000000    200005.000000  200005.000000
mean   100003.000000      0.080069      12504.441794     5.002416
min     1.000000        0.010000       5000.000000     2.000000
25%    50002.000000      0.050000      8740.000000     3.500000
50%    100003.000000      0.080000     12497.000000     5.010000
75%    150004.000000      0.120000     16264.000000     6.510000
max    200005.000000      0.150000     20000.000000     8.000000
std     57736.614632      0.040602      4337.663210     1.734485

          Date            Clicks    Impressions \
count          200005    200005.000000  200005.000000
mean  2021-07-01 23:37:44.289392896     549.774591    5507.307107
min    2021-01-01 00:00:00         100.000000    1000.000000
25%    2021-04-02 00:00:00         325.000000    3266.000000
50%    2021-07-02 00:00:00         550.000000    5518.000000
75%    2021-10-01 00:00:00         775.000000    7753.000000
max    2021-12-31 00:00:00        1000.000000   10000.000000
std                  NaN         260.019354    2596.863794

          Engagement_Score
count    200005.000000
mean         5.494673
min          1.000000
25%          3.000000
50%          5.000000
75%          8.000000
max         10.000000
std          2.872593 )

```

```

In [20]: # Get unique values for Target Audiences and Marketing Channels
unique_target_audiences = df["Target_Audience"].unique()
unique_marketing_channels = df["Channel_Used"].unique()

# Count of unique values
num_target_audiences = len(unique_target_audiences)
num_marketing_channels = len(unique_marketing_channels)

unique_target_audiences, num_target_audiences, unique_marketing_channels, num_marketing_channels

```

```

Out[20]: (array(['Men 18-24', 'Women 35-44', 'Men 25-34', 'All Ages', 'Women 25-34'],
              dtype=object),
          5,
          array(['Google Ads', 'YouTube', 'Instagram', 'Website', 'Facebook',
                'Email'], dtype=object),
          6)

```

```

In [22]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

```

```

In [27]: # Reload the dataset
file_path = (r"C:\Users\Mummy Mo\Downloads\HNG Data set\m_c_dataset_csv.csv")
df = pd.read_csv(file_path)

# Convert 'Acquisition_Cost' to numeric (removing '$' and commas if present)
df["Acquisition_Cost"] = df["Acquisition_Cost"].replace(['\$', ','], '', regex=True).astype(float)

# Set visualization style
sns.set_style("whitegrid")

# Create boxplots for detecting outliers in key numerical metrics
fig, axes = plt.subplots(1, 3, figsize=(15, 5))

sns.boxplot(y=df["Impressions"], ax=axes[0], color="skyblue")
axes[0].set_title("Impressions Distribution")


```

```
sns.boxplot(y=df["Clicks"], ax=axes[1], color="salmon")
axes[1].set_title("Clicks Distribution")

sns.boxplot(y=df["Acquisition_Cost"], ax=axes[2], color="lightgreen")
axes[2].set_title("Acquisition Cost Distribution")

plt.tight_layout()
plt.show()
```

```
<>:6: SyntaxWarning: invalid escape sequence '\$'
<>:6: SyntaxWarning: invalid escape sequence '\$'
C:\Users\Mummy Mo\AppData\Local\Temp\ipykernel_15148\2803953236.py:6: SyntaxWarning: invalid escape sequence '\$'
df["Acquisition_Cost"] = df["Acquisition_Cost"].replace(['\$'], '', regex=True).astype(float)
```



```
In [29]: # Function to detect outliers using IQR method
def detect_outliers(data, column):
    Q1 = data[column].quantile(0.25)
    Q3 = data[column].quantile(0.75)
    IQR = Q3 - Q1

    # Define lower and upper bounds for outliers
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Filter outliers
    outliers = data[(data[column] < lower_bound) | (data[column] > upper_bound)]

    return outliers

# Identify outliers for each key numerical field
outliers_impressions = detect_outliers(df, "Impressions")
outliers_clicks = detect_outliers(df, "Clicks")
outliers_spend = detect_outliers(df, "Acquisition_Cost")

# Display counts of outliers detected
outlier_counts = {
    "Impressions Outliers": len(outliers_impressions),
    "Clicks Outliers": len(outliers_clicks),
    "Acquisition Cost Outliers": len(outliers_spend)
}

outlier_counts
```

```
Out[29]: {'Impressions Outliers': 0,
          'Clicks Outliers': 0,
          'Acquisition Cost Outliers': 0}
```

```
In [10]: import pandas as pd
df = pd.read_csv(r"C:\Users\Mummy Mo\Downloads\HNG Data set\m_c_dataset_csv.csv") # Or another method to create df

# Calculate Click-Through Rate (CTR) and Cost Per Click (CPC)
df['CTR'] = df['Clicks'] / df['Impressions']
df['CPC'] = df['Acquisition_Cost'] / df['Clicks']
df['Conversion_Rate'] = df['Conversion_Rate']

# Summarize key campaign metrics
campaign_metrics = df[["CTR", "CPC", "Conversion_Rate"]].describe()

# Summary statistics for CTR and CPC
ctr_cpc_stats = df[["CTR", 'CPC', 'Conversion_Rate']].describe()

# Check for outliers in CTR and CPC using IQR method
Q1 = df[['CTR', 'CPC', 'Conversion_Rate']].quantile(0.25)
Q3 = df[['CTR', 'CPC', 'Conversion_Rate']].quantile(0.75)
IQR = Q3 - Q1
```

```

outliers = ((df[['CTR', 'CPC', 'Conversion_Rate']] < (Q1 - 1.5 * IQR)) | (df[['CTR', 'CPC', 'Conversion_Rate']]
ctr_cpc_stats, outliers

```

```

Out[10]: (
      count  200005.000000  200005.000000  200005.000000
      mean    0.140405    32.008319    0.080069
      std     0.130880    26.925841    0.040602
      min     0.010054     5.021084    0.010000
      25%     0.058606    15.092037    0.050000
      50%     0.099790    22.773973    0.080000
      75%     0.169698    38.598253    0.120000
      max     0.992024    199.960000    0.150000,
      CTR      16141
      CPC      16045
      Conversion_Rate  0
      dtype: int64)

```

```

In [39]: # Calculate Click-Through Rate (CTR) = (Clicks / Impressions) * 100
df["CTR"] = (df["Clicks"] / df["Impressions"]) * 100

# Calculate Cost Per Click (CPC) = Acquisition Cost / Clicks
df["CPC"] = df["Acquisition_Cost"] / df["Clicks"]

# Group by marketing channel and calculate average ROI, CTR, and CPC
channel_performance = df.groupby("Channel_Used")[["ROI", "CTR", "CPC", "Conversion_Rate"]].mean().reset_index()

# Sort by highest ROI
channel_performance = channel_performance.sort_values(by="ROI", ascending=False)

# Display the summary table
channel_performance

```

```

Out[39]:
   Channel_Used  ROI      CTR      CPC  Conversion_Rate
1    Facebook  5.018672  14.049724  32.129366         0.079990
4     Website  5.014114  14.096941  31.779148         0.080182
2   Google Ads  5.003126  13.918943  32.308459         0.080181
0      Email  4.996487  14.054269  31.881471         0.080282
5    YouTube  4.993720  14.119755  31.872904         0.079890
3   Instagram  4.988706  14.003691  32.080786         0.079886

```

```

In [14]: # Calculate total Clicks, Impressions, and Acquisition Cost by Marketing Channel
channel_stats = df.groupby("Channel_Used").agg(
    Total_Clicks=("Clicks", "sum"),
    Total_Impressions=("Impressions", "sum"),
    Total_Acquisition_Cost=("Acquisition_Cost", "sum")
).reset_index()

# Compute CTR and CPC
channel_stats["CTR (%)"] = (channel_stats["Total_Clicks"] / channel_stats["Total_Impressions"]) * 100
channel_stats["CPC"] = channel_stats["Total_Acquisition_Cost"] / channel_stats["Total_Clicks"]

# Display the results
channel_stats

```

```

Out[14]:
   Channel_Used  Total_Clicks  Total_Impressions  Total_Acquisition_Cost  CTR (%)  CPC
0      Email      18493963      184801107      420874104      10.007496  22.757378
1    Facebook      18038175      180662496      410603426      9.984460  22.763025
2   Google Ads      18342589      185020154      418944514      9.913833  22.839988
3   Instagram      18316654      183738455      417124850      9.968873  22.772983
4     Website      18415351      183815901      416606897      10.018367  22.622805
5    YouTube      18350935      183450845      416797090      10.003189  22.712581

```

```

In [24]: # Re-import NumPy to avoid NameError
import numpy as np

# Create a new figure with two subplots
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# CTR Bar Chart with Better Spacing
bars = axes[0].bar(channel_stats["Channel_Used"], channel_stats["CTR (%)"], color="royalblue", alpha=0.8, width=0.8)
axes[0].set_title("Click-Through Rate (CTR) by Marketing Channel", fontsize=14, fontweight="bold")

```

```

axes[0].set_xlabel("Marketing Channel", fontsize=12)
axes[0].set_ylabel("CTR (%)", fontsize=12)

# Properly set x-axis ticks and labels
axes[0].set_xticks(range(len(channel_stats["Channel_Used"])))
axes[0].set_xticklabels(channel_stats["Channel_Used"], rotation=30, ha="right")

# Adjust CTR Y-axis to create more spacing
ctr_min = np.floor(channel_stats["CTR (%)"].min()) - 0.5 # Extra space at bottom
ctr_max = np.ceil(channel_stats["CTR (%)"].max()) + 0.5 # Extra space at top
ctr_ticks = np.arange(ctr_min, ctr_max, 0.5) # Keep 0.5% increments
axes[0].set_yticks(ctr_ticks)
axes[0].set_ylim(ctr_min, ctr_max)

# Move labels higher to avoid overlap
for bar in bars:
    height = bar.get_height()
    axes[0].text(bar.get_x() + bar.get_width() / 2, height + 0.3, f"{height:.2f}%", ha="center", fontsize=10, color="black")

# CPC Line Chart
sns.lineplot(x="Channel_Used", y="CPC", data=channel_stats, ax=axes[1], marker="o", color="red", linewidth=2)
axes[1].set_title("Cost Per Click (CPC) by Marketing Channel", fontsize=14, fontweight="bold")
axes[1].set_xlabel("Marketing Channel", fontsize=12)
axes[1].set_ylabel("CPC ($)", fontsize=12)


# Properly set x-axis ticks and labels for CPC chart
axes[1].set_xticks(range(len(channel_stats["Channel_Used"])))
axes[1].set_xticklabels(channel_stats["Channel_Used"], rotation=30, ha="right")

# Add data labels on CPC points
for index, row in channel_stats.iterrows():
    axes[1].text(index, row["CPC"] + 0.05, f"${row['CPC']:.2f}", ha="center", fontsize=10, color="red")

# Adjust layout for better spacing
plt.tight_layout()

# Show the improved charts without errors
plt.show()

```



```

In [26]: # Calculate the average ROI for each marketing channel
roi_table = df.groupby("Channel_Used", as_index=False)["ROI"].mean().sort_values(by="ROI", ascending=False)

# Create a heatmap visualization
plt.figure(figsize=(8, 6))
roi_pivot = roi_table.set_index("Channel_Used") # Set marketing channel as index


# Generate the heatmap
sns.heatmap(roi_pivot, annot=True, fmt=".2f", cmap="coolwarm", linewidths=0.5, linecolor="black")

# Chart labels and title
plt.title("ROI by Marketing Channel (Heatmap)", fontsize=14, fontweight="bold")
plt.xlabel("")
plt.ylabel("Marketing Channel")

# Show the table and heatmap
plt.show()

# Display the ROI table
roi_table

```



Out[26]:

	Channel_Used	ROI
1	Facebook	5.018672
4	Website	5.014114
2	Google Ads	5.003126
0	Email	4.996487
5	YouTube	4.993720
3	Instagram	4.988706

```
In [36]: import pandas as pd

# Calculate CTR (Click-Through Rate)
df["CTR (%)"] = (df["Clicks"] / df["Impressions"]) * 100

# Create a table summarizing campaign performance by location
location_table = df.groupby("Location", as_index=False).agg({
    "ROI": "mean",
    "CTR (%)": "mean",
    "CPC": "mean",
    "Conversion_Rate": "mean"
}).sort_values(by="ROI", ascending=False) # Sort by highest ROI

# Display the table
print(location_table)
```

	Location	ROI	CTR (%)	CPC	Conversion Rate
3	Miami	5.012282	14.024957	32.152425	0.080047
2	Los Angeles	5.010876	14.067175	32.078189	0.080013
1	Houston	5.007174	14.059033	31.829355	0.079949
0	Chicago	5.001555	14.045011	32.055853	0.080131
4	New York	4.980185	14.006619	31.923819	0.080203

```
In [42]: import seaborn as sns
import matplotlib.pyplot as plt

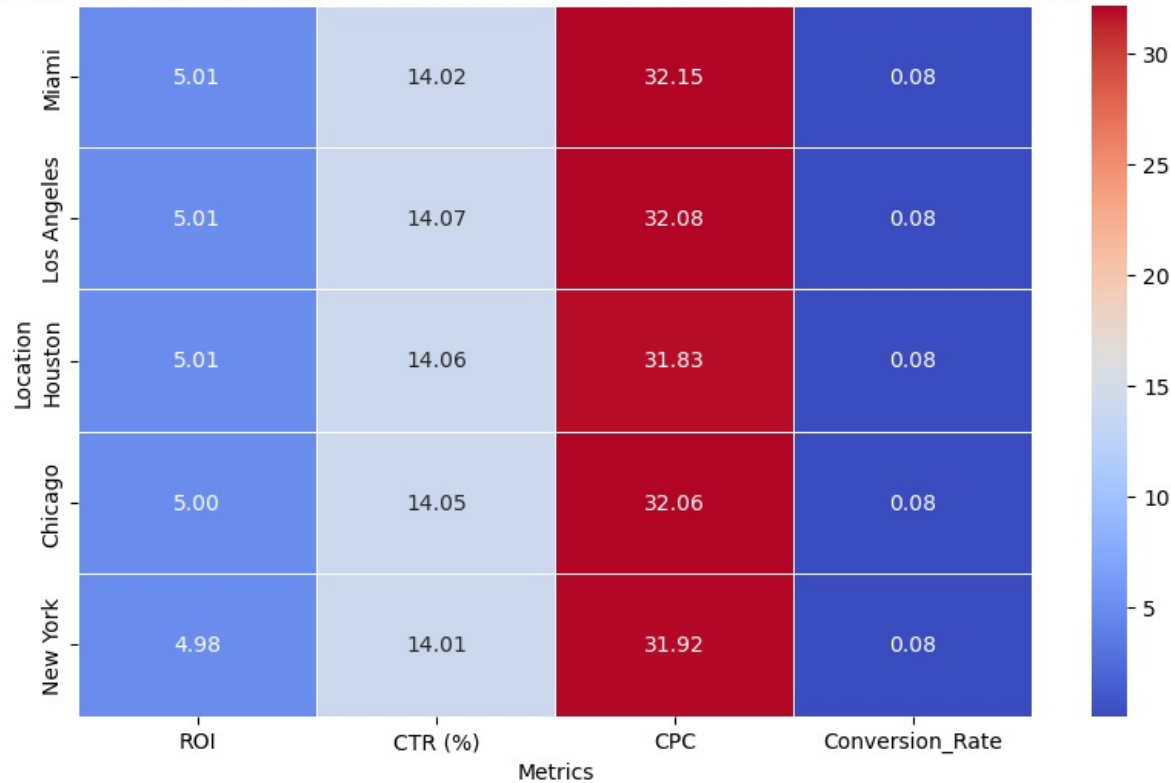
# Pivot data for heatmap
heatmap_data = location_table.set_index("Location")

# Create heatmap
plt.figure(figsize=(10, 6))
sns.heatmap(heatmap_data, annot=True, fmt=".2f", cmap="coolwarm", linewidths=0.5)

# Labels
plt.title("Location-Based Campaign Performance (ROI, CTR, Conversion Rate)", fontsize=14, fontweight="bold")
plt.xlabel("Metrics")
plt.ylabel("Location")

plt.show()
```

Location-Based Campaign Performance (ROI, CTR, Conversion Rate)




```
In [44]: import matplotlib.pyplot as plt
import seaborn as sns

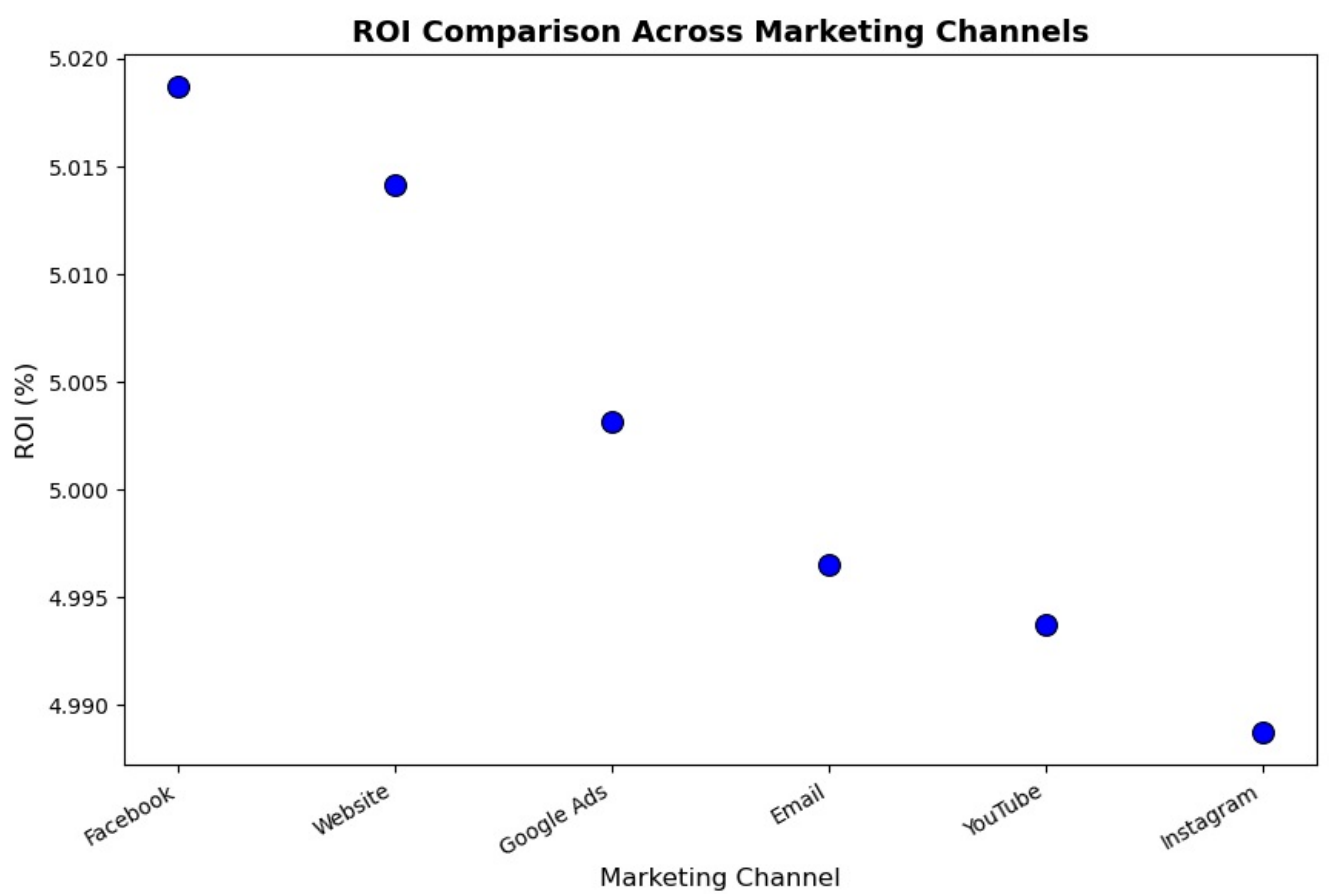
# Create scatter plot for ROI by Marketing Channel
plt.figure(figsize=(10, 6))
sns.scatterplot(data=roi_table, x="Channel_Used", y="ROI", s=100, color="blue", edgecolor="black")

# Add labels for each point
for index, row in roi_table.iterrows():
    plt.text(row["Channel_Used"], row["ROI"] + 0.1, f"{row['ROI']:.2f}%", ha="center", fontsize=10, color="black")

# Labels and title
plt.title("ROI Comparison Across Marketing Channels", fontsize=14, fontweight="bold")
plt.xlabel("Marketing Channel", fontsize=12)
plt.ylabel("ROI (%)", fontsize=12)
plt.xticks(rotation=30, ha="right")

# Show plot
plt.show()
```





In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js